



# High Dimensional Models

## Time-Varying Graphical Lasso

Andrew Boomer & Jacob Pichelmann

Toulouse School of Economics  
M2 EEE

March 17, 2021

# Overview

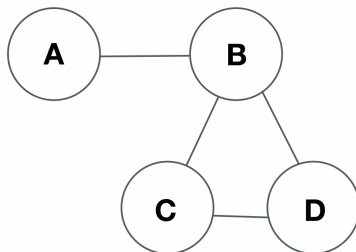
- 1 Introduction to Graphical Models
  - ▶ Important Properties
  - ▶ Interpretation
- 2 Gaussian Graphical Model
- 3 Time Varying Graphical Lasso (TGLV)
  - ▶ Altered Optimisation Problem
  - ▶ ADMM
- 4 Practical Application of TGVL
  - ▶ Comparison to Static Graphical Lasso
  - ▶ Changing the penalty function

# Graphical Models

- Graphical models offer a way to encode conditional dependencies between  $p$  random variables  $X_1, \dots, X_p$  by a graph  $g$
- A graph consists of a vertex set  $V = \{1, 2, \dots, p\}$  and an edge set  $E \subset V \times V$
- We focus on undirected graphical models, i.e. no distinction between an edge  $(s, t) \in E$  and the edge  $(t, s)$ .

Consider the following example:

**Figure:** Undirected Graphical Model



# Factorization Property

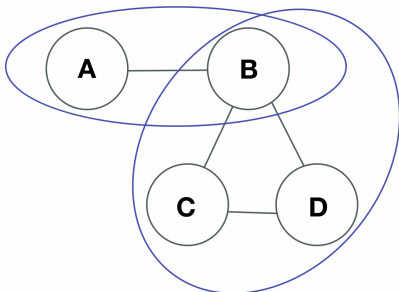
A graph clique  $C \subseteq V$  is a fully-connected subset of the vertex set, i.e.  $(s, t) \in E \forall s, t \in C$ . ( ?, ? )

$$\mathbb{P}(A, B, C, D) \propto \phi(A, B)\phi(B, C, D)$$

$$\mathbb{P}(X) = \frac{1}{Z} \prod_{c \in C} \phi_c(x_c)$$

where  $Z = \sum_{x \in X^p} \prod_{c \in C} \phi_c(x_c)$ .

Figure: Maximal Cliques

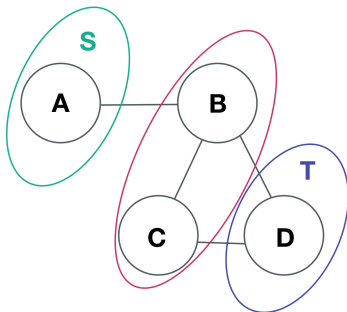


# Markov Property

Any two subsets  $S$  and  $T$  are conditionally independent given a separating subset  $Y$ . A random vector  $X$  is Markov with respect to  $g$  if

$$X_S \perp\!\!\!\perp X_T | X_Y \text{ for all cut sets } S \subset V.$$

Figure: Separating Set:  $\{B, C\}$



# Equivalence of Properties

- Hammersley-Clifford theorem:

For any strictly positive distribution the distribution of  $X$  factorizes according to the graph  $g$  if and only if the random vector  $X$  is Markov with respect to the graph. (?, ?)

# Gaussian Graphical Model

$X$  follows a Gaussian distribution:

$$X \sim \mathcal{N}(\mu, \Sigma)$$

If  $\Sigma$  is positive definite, distribution has density on  $\mathbb{R}^p$

$$f(x \mid \mu, \Sigma) = (2\pi)^{-p/2} (\det \Theta)^{1/2} e^{-(x-\mu)^T \Theta (x-\mu)/2}$$

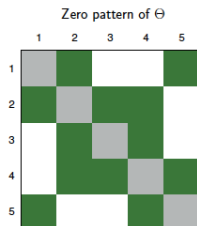
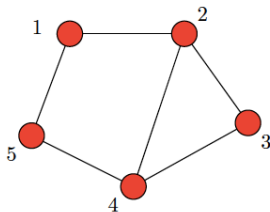
where  $\Theta = \Sigma^{-1}$  is the **Precision matrix** of the distribution.

$$\text{Empirical covariance } S = \frac{1}{n-1} \sum_{i=1}^n (x_i - \mu)(x_i - \mu)'$$



# Gaussian Graphical Model

We can represent a multivariate Gaussian distribution as a graphical model. Whenever  $X$  factorizes according to the graph  $g$  we must have  $\Theta_{st} = 0$  for any pair  $(s, t) \notin E$ . This gives a correspondence between the zero pattern of  $\Theta$  and the edge structure of  $g$ .



# Estimating the graph structure $\Leftrightarrow \Theta$

- Suppose  $X$  denotes samples from a multivariate Gaussian distribution with  $\mu = 0$  and precision matrix  $\Theta \in \mathbb{R}^{p \times p}$
- We can write the log-likelihood of the multivariate Gaussian as

$$\mathcal{L}(\Theta; X) = \frac{1}{N} \sum_{i=1}^N \log \mathbb{P}_{\Theta}(x_i) = \log \det \Theta - \text{trace}(S\Theta)$$

- So why not just estimate by MLE to obtain  $\hat{\Theta}_{\text{ML}}$  ?
  - 1 A sparse graph increases interpretability, prevents overfitting.
  - 2 In real world applications often times  $p > N$ , then MLE solution does not exist.

# $\ell_1$ Norm Regularisation

Sparsity can be achieved by adding a penalty term to the optimisation problem. Using the  $\ell_1$  norm yields the familiar lasso estimator.

$$\hat{\Theta} = \operatorname{argmin}_{\Theta \geq 0} \left( \operatorname{tr}(\mathbf{S}\Theta) - \log \det(\Theta) + \lambda \|\Theta\|_{\text{od},1} \right)$$

where  $\|\Theta\|_{\text{od},1}$  is the  $\ell_1$ -norm of the off-diagonal entries of  $\Theta$ .

# Challenge: The Network Structure Can Change Over Time

In many real world settings (e.g. financial markets) the structure of the complex system changes over time.

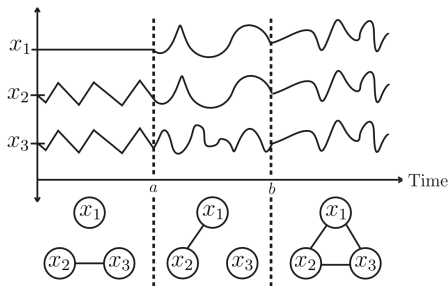


Figure: Example of Changing Network Structure (?, ?)

# Solution: Optimization on a Chain Graph (TVGL)

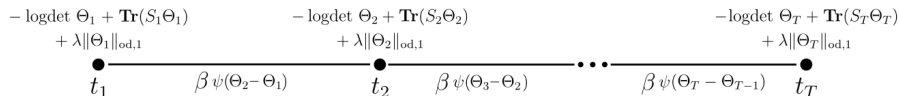


Figure: (?, ?)

The optimization problem becomes

$$\underset{\Theta \in \mathcal{S}_{++}^p}{\text{minimize}} \quad \sum_{i=1}^T -l_i(\Theta_i) + \lambda \|\Theta_i\|_{\text{od},1} + \beta \sum_{i=2}^T \psi(\Theta_i - \Theta_{i-1})$$

where  $\beta$  determines how strongly correlated neighboring covariance estimations should be. A small  $\beta$  will lead to  $\theta$ 's which fluctuate from estimate-to-estimate, whereas large  $\beta$ 's lead to smoother estimates over time.

# Choice of $\psi$

- $\psi$  allows to enforce different behaviors in the evolution of the network structure
- Expectations how the underlying network may change over time can be encoded into  $\psi$

Options:

- **Global restructuring** -  $\psi(X) = \sum_j \|[X]_j\|_2$
- **Smoothly varying over time** -  $\psi(X) = \sum_{i,j} X_{i,j}^2$
- **Perturbed node** -  $\psi(X) = \min_{V: V+V^T=X} \sum_j \|[V]_j\|_2$

# Optimization Algorithm: ADMM

- The authors use ADMM (alternating direction method of multipliers) to solve the TVGL optimization problem.
- ADMM is a general optimization technique that can be used on any convex optimization problem.
- ADMM has a couple main advantages compared to standard gradient descent based methods: (1) Can be applied to nonsmooth functions, (2) Can be distributed across multiple independent machines
- To put ADMM into context, we show how it can be used to solve a generic optimization problem

# Optimization Algorithm: ADMM

## General Example

We can take the generic minimization problem

$$\underset{x}{\operatorname{argmin}} f(x) \quad \text{s.t. } x \in C$$

And separate it into two functions,  $f$  and  $g$ , where  $g$  is the indicator of  $C$

$$\underset{x}{\operatorname{argmin}} f(x) + g(z) \quad \text{s.t. } x - z = 0$$

The variable  $z$  is known as a consensus variable, and the constraint ensures final convergence between  $x$  and  $z$



# Optimization Algorithm: ADMM

## Proximal Operators/Proximal Gradient Descent

The generality of the ADMM optimization technique relies on the method of proximal gradient descent. Proximal gradient descent makes use of proximal operators, defined as:

$$\text{prox}_{\lambda f}(v) = \underset{x}{\operatorname{argmin}} \left( f(x) + (1/2\lambda) \|x - v\|_2^2 \right)$$

The ADMM iteration based update method is:

$$x^{k+1} := \underset{x}{\operatorname{argmin}} \left( f(x) + (\rho/2) \|x - z^k + u^k\|_2^2 \right)$$

$$z^{k+1} := \Pi_C (x^{k+1} + u^k)$$

$$u^{k+1} := u^k + x^{k+1} - z^{k+1}$$

Iterations stop when  $u^k \rightarrow u^{k+1}$  ( $x - z = 0$  constraint satisfied)

# GIASSO vs. TVGL

# Importance of $\psi$

- Choice of  $\psi$  relies on knowledge about network behavior
- No a priori decision possible
- $\psi$  is fixed over time

We illustrate the importance of the choice of  $\psi$  by replicating the author's case studies with different penalty functions.

# Changing $\psi$

# References